



3,971 responses reviewed · ClinMAP-VOI v0 benchmark	17 models scored · hosted multi-provider eval	κ 0.84 blind protocol QC agreement · QA audit pass	850+ contract reviews · 94% adopted without revision
---	---	--	--

Licensed dentist and Sciences Po Global Health MPP who moved into AI evaluation because the clinic and the model fail the same way: the dangerous answer is rarely the wrong one, it is the confident one given too early. I reviewed 850+ medical AI responses under contract, then designed and shipped ClinMAP-VOI v0, a hosted benchmark built around that exact failure mode. Every claim on this page is inspectable in the repository.

EVALUATION SYSTEMS

- Metamorphic probe & benchmark design
- Eval spec authoring · schemas · claim boundaries
- Hosted multi-provider model collection
- Review queue design · policy labels · dimension scores
- Metrics pipelines · Wilson CIs · kappa statistics
- QA gates · blind QC · holdout methodology
- Reproducibility · manifests · SHA256 provenance

HUMAN DATA QUALITY

- Rubric-grounded labeling, not vibes scoring
- RLHF preference ranking · rationale discipline
- Adjudication workflow · disagreement records
- Red-flag taxonomy design
- Reviewer calibration · throughput with consistency

CLINICAL FOUNDATION

- Licensed dentist · 5 years of practice
- Pharmacovigilance · adverse event reporting
- WHO / NICE / CDC guideline anchoring
- BHBIA legal & ethical guidelines · EDC data entry

TOOLS

- Python · pandas · JSONL pipelines · git · make
- Excel/Sheets · Tableau · VOSviewer · PubMed

LANGUAGES

- Arabic (native) · English (C2) · French (B2)

EDUCATION

MPP, Global Health

Sciences Po Paris · 2023–2025
QS 2026 Politics: #1 in the EU, #3 worldwide

BDS, Hons · GPA 4.0

Ain Shams Univ. · 2012–2017

Licensed Dentist, No. 57530

Egyptian Dental Syndicate · verify: eds.eg/en/subscriptions

AI EVALUATION & BENCHMARK WORK

ClinMAP-VOI v0 — Benchmark Producer & Primary Domain Reviewer

2026

Self-directed hosted benchmark · github.com/TarekEtman/clinmap

- Designed the framework: 40 synthetic clinical decision families, 320 prompt variants, 280 metamorphic relations with oracle labels, built to test escalation timing and value-of-information judgment.
- Ran hosted collection across 17 models, then personally reviewed all 3,971 responses: policy labels, six dimension scores, evidence spans, adjudication.
- Built the QA layer instead of trusting myself: blind protocol QC (κ 0.84 vs primary), independent external holdout panel on 720 items, relation integrity 1.00, overall audit pass.
- Shipped the evidence pack: Wilson CIs, discrimination analysis, failure atlas, worked disagreement vignettes. Reproducible end-to-end with make.
- Scope stated plainly in every public artifact: synthetic probes, not clinical validation.

Medical AI Evaluator

Dec 2025 – May 2026

Outlier AI / Scale AI

- Reviewed 850+ medical AI responses spanning symptom triage, medication questions, and patient education. 94% of rationales adopted without revision; 31 pulled into rubric calibration sessions as reference standards.
- Built a personal red-flag taxonomy covering missed escalation timing, false reassurance, and unstated drug interactions; the project lead folded a version of it into the shared team rubric.
- Wrote test prompts deliberately short on history, the way real patients write them. That is where the model's worst failures showed up.

Generalist AI Trainer & LLM Evaluator

May 2025 – Dec 2025

DataAnnotation.tech

- Ranked 1,600+ general LLM outputs across reasoning, dialogue, and instruction-following, graded on what an answer assumed, not how well it read.
- Flagged a pattern that held across projects: the weaker response usually answered an easier, unstated version of the question.

SUPPORTING HEALTH & EVIDENCE WORK

Political & Scientific Officer — Pasteur Network Assoc., Institut Pasteur

Dec 2024 – May 2025

Mapped a 466-person global flavivirus expert network in Python and VOSviewer, turning scattered outbreak-response contacts into a searchable capability gap tool.

Research Analyst, AI Workflow Integration — Oracle Life Sciences

Dec 2023 – Nov 2024

Pressure-tested an Oracle Digital Assistant workflow across 3 evidence-generation projects; “sounds right” and “is traceable” kept turning out to be different bars.

Dental Practitioner & Clinic Manager — Etman Dental Centre, Cairo

Dec 2018 – Jul 2023

Five years of full clinical practice: diagnosis, treatment planning, and the actual conversation with a worried patient. Digitised records, cut overhead ~20%. IFRC COVID-19 volunteer (2020–21): trained 500+ people in risk communication.



the feather holds